

Kubernetes installation (Single node cluster)

Machine: 1 master & 2 worker

Master-node ip : 192.168.1.42

Worker-node1 ip : 192.168.1.43

Worker-node2 ip : 192.168.1.44

Host Entries: on each machine

192.168.1.42 master-node

192.168.1.43 worker-node1

192.168.1.44 worker-node2

Install podman in each machine

```
# yum install podman* -y
```

Install Cri-o in each machine

```
# curl -L -o /etc/yum.repos.d/devel:kubic:libcontainers:stable.repo
```

https://download.opensuse.org/repositories/devel:/kubic:/libcontainers:/stable/CentOS_9_Stream/devel:kubic:libcontainers:stable.repo

```
# curl -L -o /etc/yum.repos.d/devel:kubic:libcontainers:stable:cri-o:$VERSION.repo
```

https://download.opensuse.org/repositories/devel:/kubic:/libcontainers:/stable:/cri-o:/1.28/CentOS_9_Stream/devel:kubic:libcontainers:stable:cri-o:1.28.repo

```
# yum install cri-o
```

```
# systemctl status cri-o
```

```
# systemctl start cri-o
```

Add repo of kubernetes in each machine

```
export KUBE_VERSION=1.28 && \  
cat <<EOF | tee /etc/yum.repos.d/kubernetes.repo  
[kubernetes]  
name=Kubernetes
```

```
baseurl=https://pkgs.k8s.io/core:/stable:/v${KUBE_VERSION}/rpm/  
enabled=1  
gpgcheck=1  
gpgkey=https://pkgs.k8s.io/core:/stable:/v${KUBE_VERSION}/rpm/repodata/repomd.x  
ml.key  
EOF
```

Install `kubeadm`, `kubectrl`, and `kubelet` (In each machine)

```
yum install -y kubeadm kubectrl kubelet
```

```
systemctl enable --now kubelet
```

Master node

`#kubeadm init --apiserver-advertise-address=192.168.1.42/24`

`--pod-network-cidr=10.244.0.0/16`

-----OUTPUT-----

```
[root@master-node ~]# kubeadm init --apiserver-advertise-address=192.168.1.42  
--pod-network-cidr=10.244.0.0/16  
I0309 17:57:37.459333 10186 version.go:256] remote version is much newer: v1.29.2;  
falling back to: stable-1.28  
[init] Using Kubernetes version: v1.28.7  
[preflight] Running pre-flight checks  
error execution phase preflight: [preflight] Some fatal errors occurred:  
    [ERROR FileContent--proc-sys-net-ipv4-ip_forward]: /proc/sys/net/ipv4/ip_forward  
contents are not set to 1  
[preflight] If you know what you are doing, you can make a check non-fatal with  
`--ignore-preflight-errors=...`  
To see the stack trace of this error execute with --v=5 or higher  
[root@master-node ~]# Vi /etc/sysctl.conf  
-bash: Vi: command not found  
[root@master-node ~]# vi /etc/sysctl.conf  
[root@master-node ~]# sysctl -p  
net.ipv4.ip_forward = 1
```

```
[root@master-node ~]# kubeadm init --apiserver-advertise-address=192.168.1.42
--pod-network-cidr=10.244.0.0/16
I0309 17:58:38.821331 10250 version.go:256] remote version is much newer: v1.29.2;
falling back to: stable-1.28
[init] Using Kubernetes version: v1.28.7
[preflight] Running pre-flight checks
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet
connection
[preflight] You can also perform this action in beforehand using 'kubeadm config images
pull'
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [kubernetes kubernetes.default
kubernetes.default.svc kubernetes.default.svc.cluster.local master-node] and IPs
[10.96.0.1 192.168.1.42]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [localhost master-node] and
IPs [192.168.1.42 127.0.0.1 ::1]
[certs] Generating "etcd/peer" certificate and key
[certs] etcd/peer serving cert is signed for DNS names [localhost master-node] and IPs
[192.168.1.42 127.0.0.1 ::1]
[certs] Generating "etcd/healthcheck-client" certificate and key
[certs] Generating "apiserver-etcd-client" certificate and key
[certs] Generating "sa" key and public key
[kubeconfig] Using kubeconfig folder "/etc/kubernetes"
[kubeconfig] Writing "admin.conf" kubeconfig file
[kubeconfig] Writing "kubelet.conf" kubeconfig file
[kubeconfig] Writing "controller-manager.conf" kubeconfig file
[kubeconfig] Writing "scheduler.conf" kubeconfig file
[etcd] Creating static Pod manifest for local etcd in "/etc/kubernetes/manifests"
[control-plane] Using manifest folder "/etc/kubernetes/manifests"
[control-plane] Creating static Pod manifest for "kube-apiserver"
[control-plane] Creating static Pod manifest for "kube-controller-manager"
```

[control-plane] Creating static Pod manifest for "kube-scheduler"

[kubelet-start] Writing kubelet environment file with flags to file
"/var/lib/kubelet/kubeadm-flags.env"

[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"

[kubelet-start] Starting the kubelet

[wait-control-plane] Waiting for the kubelet to boot up the control plane as static Pods
from directory "/etc/kubernetes/manifests". This can take up to 4m0s

[apiclient] All control plane components are healthy after 7.503427 seconds

[upload-config] Storing the configuration used in ConfigMap "kubeadm-config" in the
"kube-system" Namespace

[kubelet] Creating a ConfigMap "kubelet-config" in namespace kube-system with the
configuration for the kubelets in the cluster

[upload-certs] Skipping phase. Please see --upload-certs

[mark-control-plane] Marking the node master-node as control-plane by adding the
labels: [node-role.kubernetes.io/control-plane
node.kubernetes.io/exclude-from-external-load-balancers]

[mark-control-plane] Marking the node master-node as control-plane by adding the
taints [node-role.kubernetes.io/control-plane:NoSchedule]

[bootstrap-token] Using token: 0dvfow.sebdmxm6qwi8d6a4

[bootstrap-token] Configuring bootstrap tokens, cluster-info ConfigMap, RBAC Roles

[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to get nodes

[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to post CSRs
in order for nodes to get long term certificate credentials

[bootstrap-token] Configured RBAC rules to allow the csrapprover controller
automatically approve CSRs from a Node Bootstrap Token

[bootstrap-token] Configured RBAC rules to allow certificate rotation for all node client
certificates in the cluster

[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace

[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet
client certificate and key

[addons] Applied essential addon: CoreDNS

[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Alternatively, if you are the root user, you can run:

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:

<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 192.168.1.42:6443 --token 0dvfow.sebdmxm6qwi8d6a4 \
--discovery-token-ca-cert-hash
sha256:4cebf59a7fe8fa00f4b89130b3107e7b0e9b0f7627429d4bd29f56953028e5c5
```

Run on Master node

```
# mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

Download

```
# curl
```

```
https://raw.githubusercontent.com/projectcalico/calico/v3.26.3/manifests/canal.yaml
-O
```

```
# kubectl apply -f canal.yaml
```

If cidr not define in master initialization time , then,

Vi canal.yaml

```
# - name: CALICO_IPV4POOL_CIDR
#   value: "10.244.0.0/16"
```

Set network and save it.

Error: [root@worker-node1 ~]# kubeadm join 192.168.1.42:6443 --token
0dvfow.sebdmxm6qwi8d6a4 --discovery-token-ca-cert-hash
sha256:4cebf59a7fe8fa00f4b89130b3107e7b0e9b0f7627429d4bd29f56953028e5c5
[preflight] Running pre-flight checks
[WARNING Service-Kubelet]: kubelet service is not enabled, please run 'systemctl
enable kubelet.service'
error execution phase preflight: [preflight] Some fatal errors occurred:
[ERROR FileContent--proc-sys-net-ipv4-ip_forward]: /proc/sys/net/ipv4/ip_forward
contents are not set to 1
[preflight] If you know what you are doing, you can make a check non-fatal with
`--ignore-preflight-errors=...`
To see the stack trace of this error execute with --v=5 or higher

Solution: # vi /etc/sysctl.conf
Add line: net.ipv4.ip_forward = 1

And save it.
sysctl -p

Adding working node in cluster (run on worker nodes)

```
kubeadm join 192.168.1.42:6443 --token 0dvfow.sebdmxm6qwi8d6a4 \  
--discovery-token-ca-cert-hash  
sha256:4cebf59a7fe8fa00f4b89130b3107e7b0e9b0f7627429d4bd29f56953028e5c5
```

kubectl get nodes -o wide

```
# kubectl get pods -A
```

```
#kubectl get pods -o wide -A
```